

Engineering Swarm Behaviours through (Hybrid) Aggregate Computing

Gianluca Aguzzi gianluca.aguzzi@unibo.it

Alma Mater Studiorum – Università di Bologna

08/04/2024



Outline

1 Swarm Behaviors – Overview

2 Aggregate Computing Toolchain

- ScaFi
- MacroSwarm

3 Hybrid Aggregate Computing

- Application-level
- Runtime-level

Swarm Behaviors

Definition

A *swarm behavior* a collective behavior that emerges from the interaction of a group of agents.

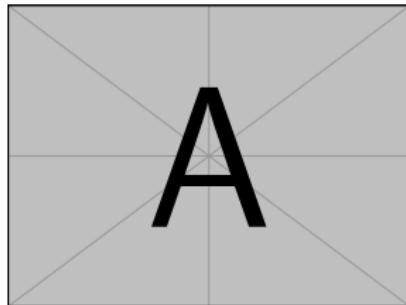
Taxonomy

From a swarm robotics perspective, swarm behaviors can be classified into three main categories [1]:

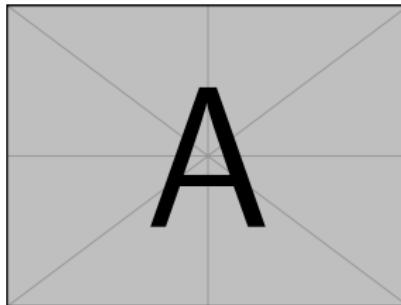
- *Spatial-organization behaviors*: the agents organize themselves in a specific spatial configuration
- *Navigation behaviors*: the agents collectively navigate in the environment following a specific strategy
- *Collective Decision-Making*: the agents collectively make decisions based on the information they have

[1] M. Brambilla, E. Ferrante, M. Birattari, *et al.*, "Swarm robotics: A review from the swarm engineering perspective," *Swarm Intell.*, vol. 7, no. 1, pp. 1–41, 2013. doi: 10.1007/S11721-012-0075-2. [Online]. Available: <https://doi.org/10.1007/s11721-012-0075-2>

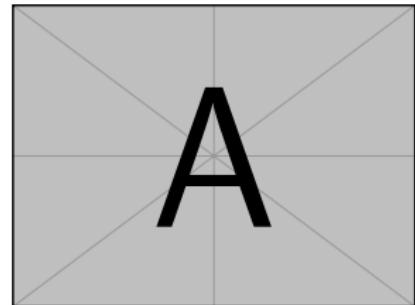
Spatial-organization behaviors



Aggregation



Dispersion

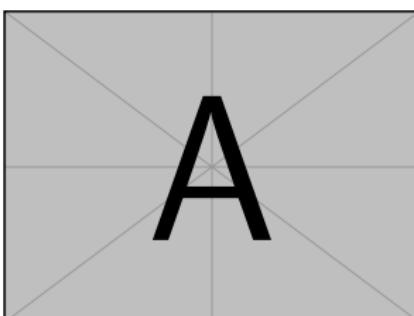


Flocking

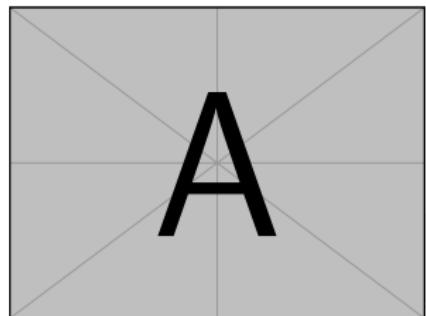
Navigation behaviors



Exploration

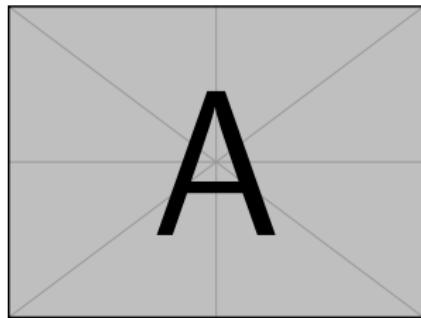


Collective Motion

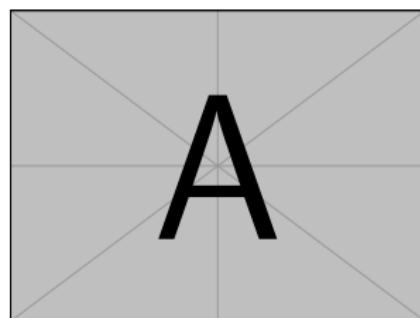


Group Transportation

Collective Decision-Making



Exploration



Collective Motion

Engineering Swarm Behaviors I

Definition

Engineering swarm behaviors is the process of designing, implementing, and deploying algorithms that enable a group of agents to exhibit a specific collective behavior.

Challenges

- *Scalability*: the collective behavior should scale with the number of agents
- *Robustness*: even in case of failures, the collective behavior should be maintained
- *Flexibility*: the collective behavior should be easily adaptable to different scenarios
- *Efficiency*: the collective behavior should be efficient in terms of energy and time



Engineering Swarm Behaviors II

Properties

- *Emergence*: the collective behavior emerges from the interaction of the agents
- *Decentralization*: the collective behavior is achieved without a central controller
- *Self-organization*: the collective behavior is achieved through local interactions
- *Homogeneity*: the agents are homogeneous and have the same capabilities
- *Limited Communication*: the agents have limited communication capabilities

State-of-the-art

- *manual design*: the collective behavior is manually designed through
 - bio-inspired algorithms, macroprogramming languages → aggregate programming
- *automatic design*: the collective behavior is automatically designed through machine learning techniques
 - reinforcement learning, evolutionary algorithms, multi-agent reinforcement learning
- *hybrid design*: a combination of manual and automatic design

Aggregate Programming – What we want

Goal

- a **minimal core** to cover the main functionality of field calculus
- a **set of libraries** that cover the main building blocks
- an **high-level API** for domain-specific applications

Challenges

- Support several kinds of devices
 - *constrained, commodity, consumer, server*
- Support several infrastructures
 - *peer-to-peer, server-based, cloud-edge*
- Expressive and simple enough to reduce the cognitive load

Aggregate programming: current implementations I

A large number of incarnations and implementations have been developed over the years, with different *trade-offs* and *goals*.

External DSL

- **Proto** [2]: first incarnation of FC
 - *good performance* → *low-level*
 - *hard to use* → *hard to maintain*
- **Protelis** [3]: an external DSL based on Xtext
 - *good ergonomics* → *high-level*
 - *extensive libraries* → *easy to use*
 - *hard to maintain*
 - *JVM-based* → *hard to port to constrained devices*

Aggregate programming: current implementations II

Internal DSL

- **FCPP** [4]: a performance-oriented internal DSL in C++
 - *good performance* → *low-level*
 - *hard to use* → *bad ergonomics*
- **ScaFi** [5]: a high-level internal DSL in Scala (current reference implementation and focus of this talk)
 - *good ergonomics* → *high-level*
 - *extensive libraries* → *easy to use*
 - *Cross-platform*
 - *Scala Native* → *hard to port to constrained devices*

Experimental

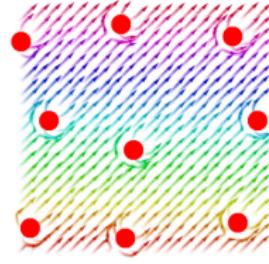
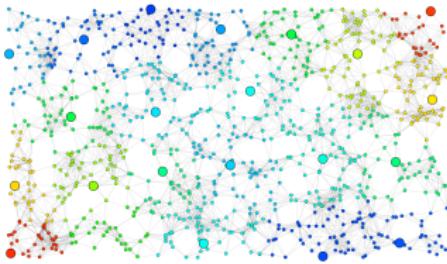
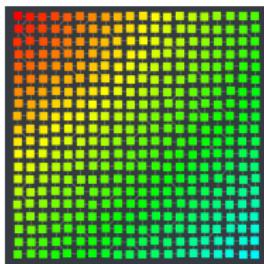
- **Collektive**^a: a high-level DSL in Kotlin leveraging the compiler plugin and the Kotlin Multiplatform
- **RuFi**^b: a minimal core in Rust (WIP)

^a<https://github.com/Collektive/collektive>

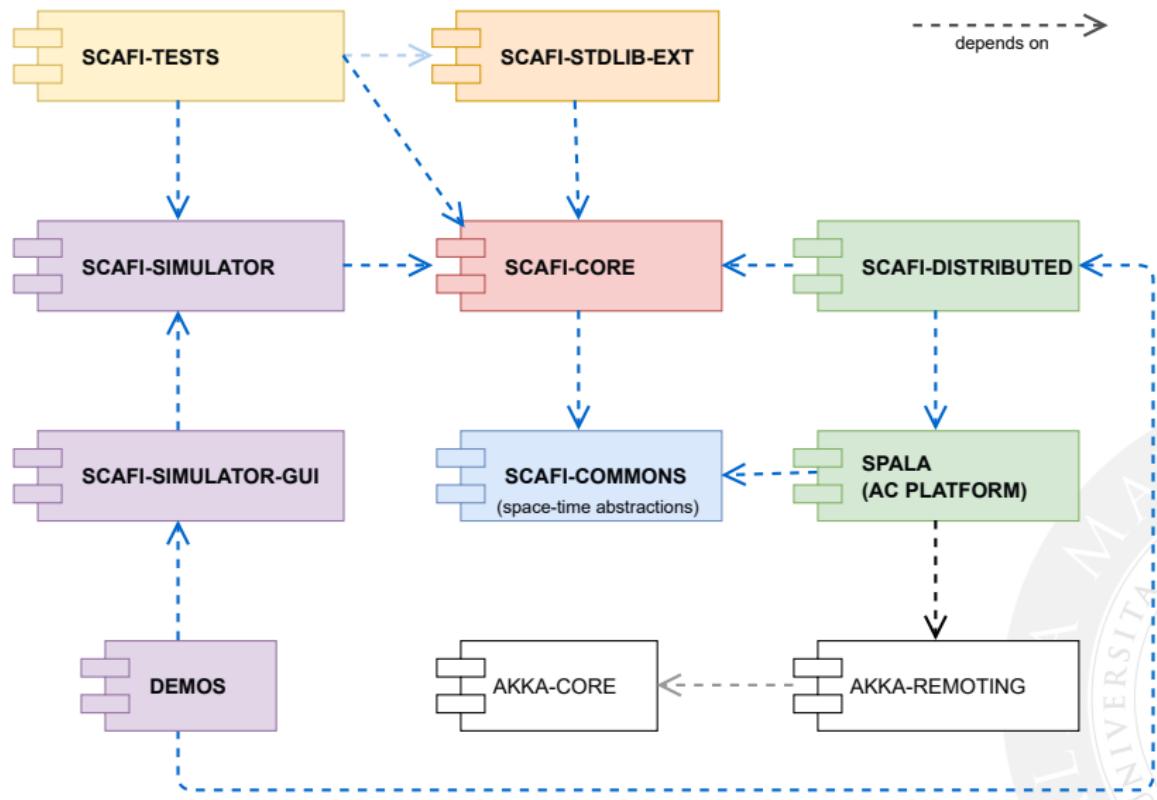
^b<https://github.com/lm98/rufi>

ScaFi (Scala Fields)

A Scala toolkit providing an *internal domain-specific* language, *libraries*, a *simulation* environment, and *runtime* support for **practical** aggregate computing systems development



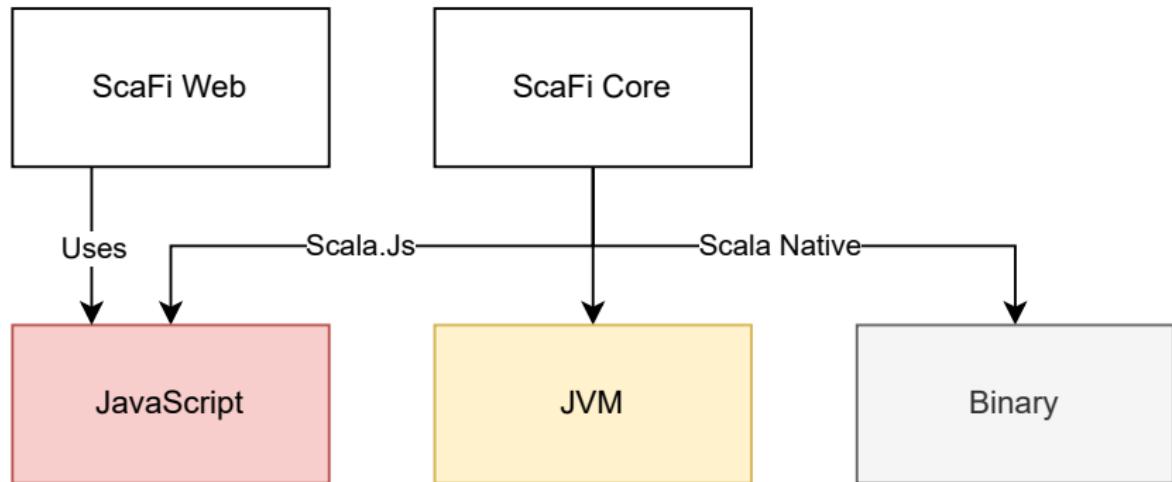
ScaFi: Organization



ScaFi: syntax as a core language/API

```
trait Constructs {  
    // the unique identifier of the local device  
    def mid(): ID  
  
    // applies fun to the previous result, or to init at the first call  
    def rep[A](init: A)(fun: (A) => A): A  
  
    // evaluation of expr at the currently-considered neighbor  
    def nbr[A](expr: => A): A  
  
    // accumulates available evaluations of expr, with acc/init monoid  
    def foldhood[A](init: => A)(acc: (A,A)=>A)(expr: => A): A  
  
    // splits computation: th where cond is true, el everywhere/time else  
    def branch[A](cond: => Boolean)(th: => A)(el: => A): A  
  
    // perception of local sensor  
    def sense[A](name: CNAME): A  
  
    // perception of neighbourhood sensor  
    def nbrvar[A](name: CNAME): A  
    ...  
}
```

ScaFi: Cross-platform



ScaFi: setup an aggregate application

```
package experiments
// STEP 1 Choose an incarnation (simulated or real)
import it.unibo.scafi.incarnations.BasicSimulationIncarnation._

// STEP 2 Define the aggregate program including the right libraries
class MyProgram extends AggregateProgram with Libs {
    // STEP 2.1 Define main logic of the program
    override def main(): Any = rep(0)(_ + 1)
}

// STEP 3 Platform/Node setup (both in simulation/real)

// STEP 3.1 in case of ScaFi simulation:
object SimulationRunner extends Launcher {
    Settings.Sim_ProgramClass = "experiments.MyAggregateProgram"
    Settings.ShowConfigPanel = true
    launch()
}
```

An aggregate computing playground – ScaFi Web!

<https://scafi.github.io/web/>

The screenshot shows the ScaFi Web interface. On the left, there are three configuration panels: 'Backend configuration' (with code editor showing Scala code for a sensor), 'Led matrix configuration' (color: red, dimension: 3x3, static: true), and 'Sensors configuration' (add sensor: sensor, false). The main area features a color gradient heatmap on a 3x3 grid, with a central red square. The top navigation bar includes 'Dark / Light', 'Settings', 'Website', 'Repository', and other icons.

```
1 //using StandardSensors, Actuation
2 val gradient = rep[Double, PositiveInfinity](distance =>
3   mux(sense[Boolean]("sensor"))(
4     0.0
5     M
6     minHoodPlus(nbr(distance) + nbrRange)
7   )
8 )
9 val maxValue = 500.0
10 val hueColor = if(gradient.isInfinite) { 0 } else { gradient / max
11 val saturation = if(gradient.isInfinite) { 0 } else { gradient / r
12 val color = hsl(hueColor, 1, 0.5)
13 (gradient.formatted("%4.2f"), ledAll to color)
```

ScaFi Blocks

ScaFi Blocks is a visual programming environment for ScaFi.

Scafi

Discover the power of the collective

Settings Website Repository

load start stop tick speed Slow Normal Fast

id neighborhood export Sensors

Show code

Aggregate Program

Output Channel Source Target Width

Sense Boolean "source"

Sense Boolean "target"

5

1 8 15 22 29 36 43
2 9 16 23 30 37 44
3 10 17 24 31 38 45
4 11 18 25 32 39 46
5 12 19 26 33 40 47
6 13 20 27 34 41 48
7 14 21 28 35 42 49

The screenshot shows the ScaFi Blocks interface. On the left, there's a sidebar with a 'Show code' button and a list of categories: Basic, Conditions, Building Blocks, User-API, Sensing, Actuation, Movement, Operations, Values, Types, and Definitions. The main workspace displays a visual program consisting of several blocks: an 'Output' block, a 'Channel' block, a 'Source' block, a 'Target' block, and a 'Width' block. The 'Source' and 'Target' blocks each have a 'Sense' block connected to them. The 'Sense' blocks are configured with 'Boolean' types and labels 'source' and 'target'. A numerical value '5' is also present. To the right of the workspace is a large network visualization showing 49 nodes (labeled 1 to 49) connected by a complex web of lines, representing a neighborhood or sensor field. Above the network are various control buttons: 'load', 'start', 'stop', 'tick', 'speed' (with options 'Slow', 'Normal', 'Fast'), 'id', 'neighborhood', 'export', and 'Sensors'. At the bottom right is a large blue play button.

MacroSwarm

What

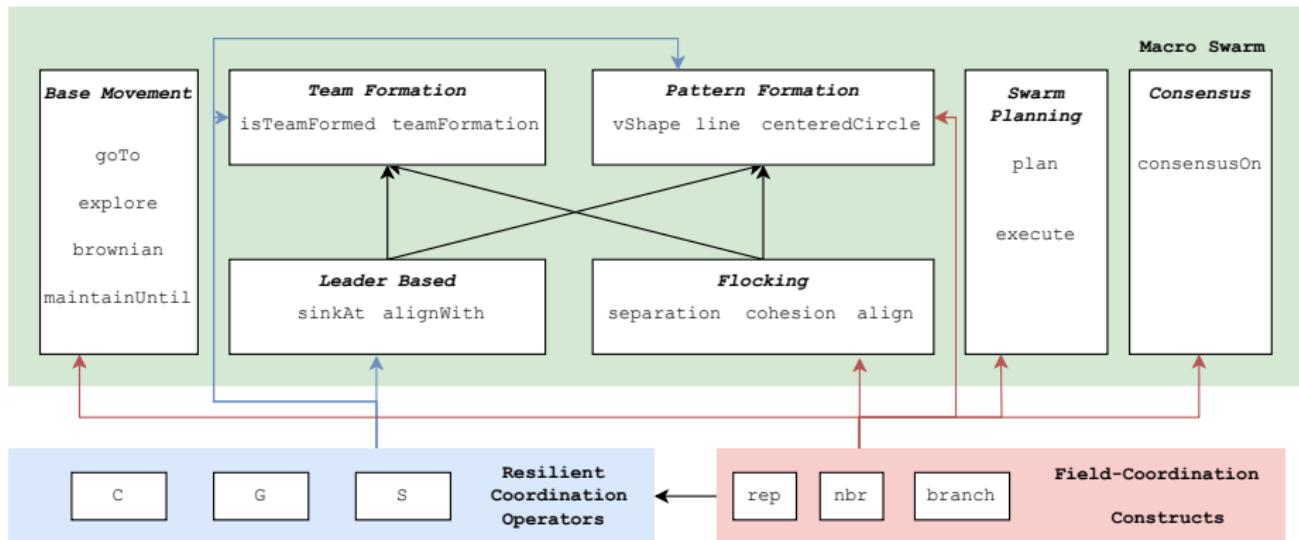
MacroSwarm^a [6] is a ScaFi library for programming swarm robotics applications.

^a<https://scafi.github.io/macro-swarm/>

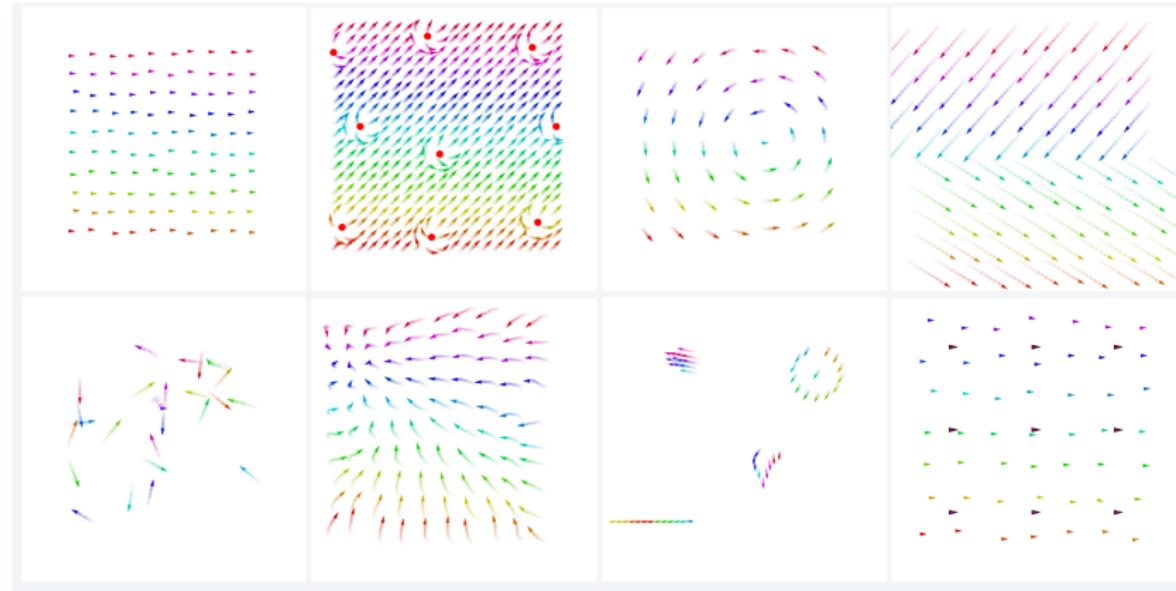
How

- A MacroSwarm program is a function from *sensing* field to *action* field
 - *sensing* field: the information perceived by the agents (e.g., the distance from a source)
 - *action* field: the action to be performed by the agents (e.g., a velocity vector)
- The action field then should be translated into the actual action of the agents based on the specific robot capabilities
 - e.g., a velocity vector can be translated into motor commands

MacroSwarm: API structure



MacroSwarm: Example



Towards Hybrid Aggregate Computing

What

Hybrid Aggregate Computing is the process of combining *manual* (aggregate programming) and *automatic* design techniques to engineer swarm behaviors.

Why

- *Manual design* is good for designing *simple* and *well-understood* behaviors
- *Automatic design* is good for designing *complex* and *unknown* behaviors
- *Hybrid design* is good for combining the best of both worlds



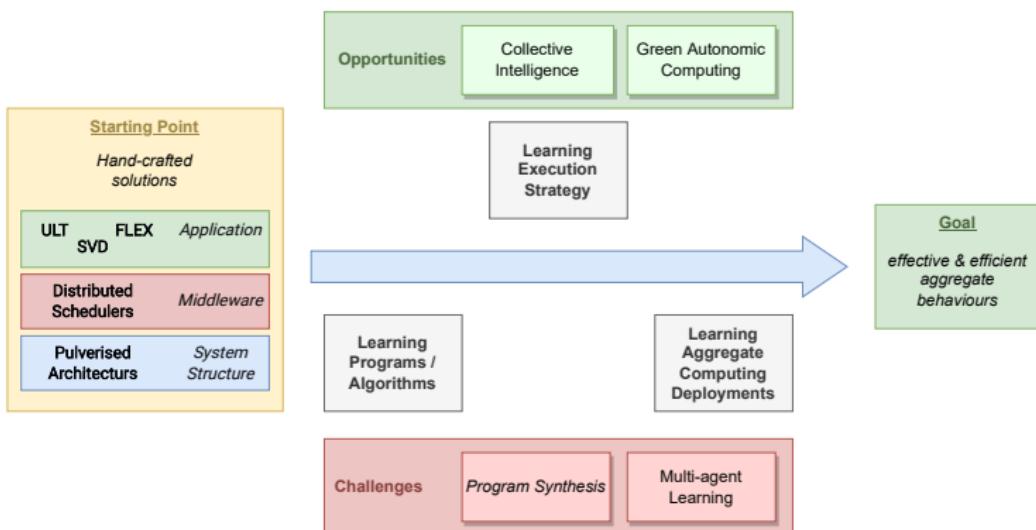
Research Agenda

Goals:

- functionality (e.g., a certain collective behaviour)
- non-functionality (e.g., energy consumption)

Means:

- system structures
- execution strategy
- algorithms



Hybrid Aggregate Computing – Execution Strategy

Distributed schedulers

- The aggregate computing model does not enforce a global-synchronization
- Distributed schedulers [7]: local decisions on the rounds of the aggregate program
- Idea: learning the best scheduling strategy directly by-doing

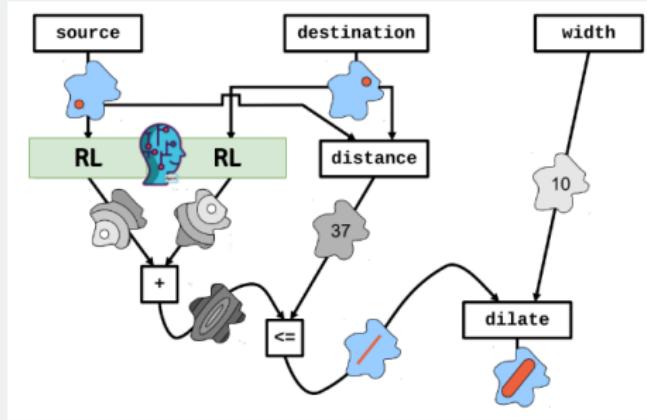


[7] G. Aguzzi, R. Casadei, and M. Viroli, "Addressing collective computations efficiency: Towards a platform-level reinforcement learning approach," in *IEEE International Conference on Autonomic Computing and Self-Organizing Systems, ACSOS 2022, Virtual, CA, USA, September 19–23, 2022*, R. Casadei, E. D. Nitto, I. Gerostathopoulos, *et al.*, Eds., IEEE, 2022, pp. 11–20. doi: 10.1109/ACSOS55765.2022.00019. [Online]. Available: <https://doi.org/10.1109/ACSOS55765.2022.00019>

Hybrid Aggregate computing – Algorithms

Collective Program Sketching

- holes: blocks of the aggregate program depending on environment dynamics
- sketching: partial specification of the aggregate program
- Idea [8]: synthesis of the holes by learning from the realistic simulation

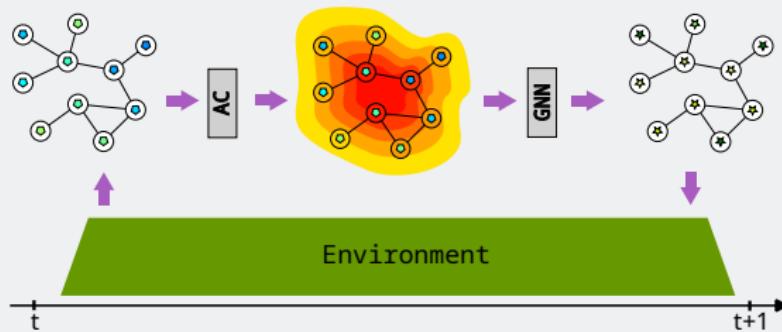


[8] G. Aguzzi, R. Casadei, and M. Viroli, "Towards reinforcement learning-based aggregate computing," in *Coordination Models and Languages - 24th IFIP WG 6.1 International Conference, COORDINATION 2022*, M. H. ter Beek and M. Sirjani, Eds., ser. Lecture Notes in Computer Science, vol. 13271, Springer, 2022, pp. 72–91. doi: 10.1007/978-3-031-08143-9_5. [Online]. Available: https://doi.org/10.1007/978-3-031-08143-9_5

Hybrid Aggregate computing – Algorithms

Field-Informed reinforcement learning

- Aggregate computing is used to inform the learning process
 - Speeding up the training & Reducing the search space [9]
- Field used as a way to produce a stygmeric representation of the environment
- Learning performed through deep reinforcement learning
- Graph-neural networks used as a policy function approximator



[9] G. Aguzzi, M. Viroli, and L. Esterle, "Field-informed reinforcement learning of collective tasks with graph neural networks," in *4th IEEE International Conference on Autonomic Computing and Self-Organizing Systems - ACSOS 2023, Toronto, 2023*

Conclusion

Takeaways

- *Swarm behaviors* are collective behaviors that emerge from the interaction of a group of agents
- *Aggregate computing* is a programming model that provides a high-level API for engineering swarm behaviors
- *ScaFi* is a Scala toolkit that provides an internal DSL, libraries, a simulation environment, and runtime support for practical aggregate computing systems development
- *MacroSwarm* is a ScaFi library for programming swarm robotics applications
- *Hybrid Aggregate Computing* is the process of combining manual and automatic design techniques to engineer swarm behaviors



Engineering Swarm Behaviours through (Hybrid) Aggregate Computing

Gianluca Aguzzi gianluca.aguzzi@unibo.it

Alma Mater Studiorum – Università di Bologna

08/04/2024



References |

- [1] M. Brambilla, E. Ferrante, M. Birattari, and M. Dorigo, "Swarm robotics: A review from the swarm engineering perspective," *Swarm Intell.*, vol. 7, no. 1, pp. 1–41, 2013. doi: 10.1007/S11721-012-0075-2. [Online]. Available: <https://doi.org/10.1007/s11721-012-0075-2>.
- [2] J. Beal and J. Bachrach, "Infrastructure for engineered emergence on sensor/actuator networks," *IEEE Intell. Syst.*, vol. 21, no. 2, pp. 10–19, 2006. doi: 10.1109/MIS.2006.29. [Online]. Available: <https://doi.org/10.1109/MIS.2006.29>.
- [3] D. Pianini, M. Viroli, and J. Beal, "Protelis: Practical aggregate programming," in *Proceedings of the 30th Annual ACM Symposium on Applied Computing, Salamanca, Spain, April 13-17, 2015*, R. L. Wainwright, J. M. Corchado, A. Bechini, and J. Hong, Eds., ACM, 2015, pp. 1846–1853. doi: 10.1145/2695664.2695913. [Online]. Available: <https://doi.org/10.1145/2695664.2695913>.
- [4] G. Audrito, "FCPP: an efficient and extensible field calculus framework," in *IEEE International Conference on Autonomic Computing and Self-Organizing Systems, ACSOS 2020, Washington, DC, USA, August 17-21, 2020*, IEEE, 2020, pp. 153–159. doi: 10.1109/ACSOS49614.2020.00037. [Online]. Available: <https://doi.org/10.1109/ACSOS49614.2020.00037>.
- [5] R. Casadei, M. Viroli, G. Aguzzi, and D. Pianini, "Scafì: A scala DSL and toolkit for aggregate programming," *SoftwareX*, vol. 20, p. 101248, 2022. doi: 10.1016/j.softx.2022.101248. [Online]. Available: <https://doi.org/10.1016/j.softx.2022.101248>.
- [6] G. Aguzzi, R. Casadei, and M. Viroli, "Macroswarm: A field-based compositional framework for swarm programming," *CoRR*, vol. abs/2401.10969, 2024. doi: 10.48550/ARXIV.2401.10969. arXiv: 2401.10969. [Online]. Available: <https://doi.org/10.48550/arXiv.2401.10969>.

References II

- [7] G. Aguzzi, R. Casadei, and M. Viroli, "Addressing collective computations efficiency: Towards a platform-level reinforcement learning approach," in *IEEE International Conference on Autonomic Computing and Self-Organizing Systems, ACSOS 2022, Virtual, CA, USA, September 19-23, 2022*, R. Casadei, E. D. Nitto, I. Gerostathopoulos, *et al.*, Eds., IEEE, 2022, pp. 11–20. doi: 10.1109/ACSOS55765.2022.00019. [Online]. Available: <https://doi.org/10.1109/ACSOS55765.2022.00019>.
- [8] G. Aguzzi, R. Casadei, and M. Viroli, "Towards reinforcement learning-based aggregate computing," in *Coordination Models and Languages - 24th IFIP WG 6.1 International Conference, COORDINATION 2022*, M. H. ter Beek and M. Sirjani, Eds., ser. Lecture Notes in Computer Science, vol. 13271, Springer, 2022, pp. 72–91. doi: 10.1007/978-3-031-08143-9_5. [Online]. Available: https://doi.org/10.1007/978-3-031-08143-9_5.
- [9] G. Aguzzi, M. Viroli, and L. Esterle, "Field-informed reinforcement learning of collective tasks with graph neural networks," in *4th IEEE International Conference on Autonomic Computing and Self-Organizing Systems - ACSOS 2023, Toronto*, 2023.